

Multimodal Benchmark Evaluation Dossier

Doc ID: EVAL-MM-2026-X7 • Target: Vision-Language Model QA & Extraction • Classification: Test Artifact

Charts: Bar, Line, Donut

Text Modality: Embedded

Evaluation Domain: Tech & Finance

This benchmark document assesses multimodal LLM capabilities across cross-modal reasoning, quantitative graph extraction, and visual element association.

Section 1: Enterprise Revenue & Profit Dynamics

In fiscal years 2025 and 2026, NexusTech recorded significant growth across international and enterprise business units. The table and chart below illustrate financial metrics across 6 operational quarters.

Figure 1: Quarterly Financial Performance (2025 - 2026)




Figure 1: NexusTech quarterly revenue (dark blue) vs. net profit (teal) in millions USD.

Period	Primary Revenue Driver	Reported Headcount	Operating Margin Status
Q1 2025 – Q2 2025	Legacy Cloud Hosting & SLA Contracts	420 FTE	Within Target (16.9% – 19.5%)
Q3 2025 – Q4 2025	Enterprise AI Assistant Tier Launch	485 FTE	Exceeded Target (20.3% – 23.6%)
Q1 2026 – Q2 2026	Autonomous Agent Platform Subscriptions	560 FTE	Historical Peak (22.2% – 24.1%)

Section 2: Latency vs. Model Quantization Scaling

Evaluating inference efficiency is critical when choosing between uncompressed 16-bit floating-point weights (FP16) and 8-bit integer quantization (INT8). Latency benchmarks were run across variable batch sizes (1 to 64).

Figure 2: Inference Latency vs Batch Size Across Precision Modes




Figure 2: End-to-end inference latency comparing FP16 baseline against INT8 quantized engine.

Key Operational Finding: The INT8 quantized model retains 99.1% of baseline benchmark accuracy while cutting processing latency by roughly 43% to 48% across all evaluated batch ranges.

Section 3: Fiscal 2026 Infrastructure Allocation

Capital distribution for AI operations across five core units has been partitioned for the 2026 fiscal cycle to ensure cluster sustainability and compliance with incoming regulations.

Figure 3: Fiscal 2026 AI Infrastructure Budget Breakdown




Figure 3: Relative budget proportion across five internal operational departments.

The total allocated budget across all categories sums to exactly \$140.0 Million USD, with Cloud Infrastructure taking priority for cluster reservation commitments.